

Data-driven recommendations for increasing reliability and validity in measures of infant word recognition

Veronica Boyce^{a,b}, Virginia A. Marchman^b, Heidi Baumgartner^b, Claire Augusta Bergey^c, Mika Braginsky^b, Federico Giovannetti^d, George Kachergis^b, Jess Mankewitz^e, Stephan Meylan^f, Ben Prystawski^b, Robert Z. Sparks^{b,g}, Adrian Steffan^h, Alvin Wei Ming Tan^b, Michael C. Frank^b, & Martin Zetterstenⁱ

^a Department of Brain and Cognitive Sciences, Massachusetts Institute of Technology, Cambridge, MA, USA

^b Department of Psychology, Stanford University, Stanford, CA, USA

^c Department of Linguistics, Stanford University, Stanford, CA, USA

^d Unidad de Neurobiología Aplicada (UNA, CEMIC-CONICET), Buenos Aires, Argentina

^e Department of Psychology, University of Wisconsin, Madison, WI, USA

^f Department of Linguistics, University of California, Berkeley, CA, USA

^g Department of Psychology, Harvard University, Cambridge, MA, USA

^h Department of Psychology, Ludwig Maximilian University, Munich, Germany

ⁱ Department of Cognitive Science, University of California, San Diego, CA, USA

Author Note

Correspondence concerning this article should be addressed to Veronica Boyce. E-mail: vboyce@mit.edu. We thank Kyle MacDonald and Tian Linger Xu for their critical role in establishing Peekbank and in the beginning stages of this project.

Author contributions: conceptualization – VAM, MCF, MZ; data curation – all authors; formal analysis – VB, GK, AWMT, MCF, MZ; funding acquisition – VAM, MCF, MZ; investigation – VB, CAB, GK, AWMT, MCF, MZ; methodology – VB, VAM, GK, AWMT, MCF, MZ; project administration – VAM, MCF, MZ; software – CAB, MB, GK, JM, SM, BP, AS, AWMT, MCF, MZ; supervision – VAM, MCF, MZ; validation – VB, HB, JM, BP, RZS, AWMT, MCF, MZ; visualization – VB, MCF, MZ; writing – original draft – VB, MCF, MZ; writing – reviewing and editing – all authors.

Funding: MZ was supported by a CRISP grant from the Jacobs foundation. VB was supported by the National Science Foundation under Grant No. SMA-2507597.

AI Disclosure statement: Generative AI was used to do code refactoring, code review, and proofreading.

Abstract

Children’s early word recognition is both an important phenomenon in its own right and a key index of early language more broadly. A central paradigm for measuring word recognition in young children is the looking-while-listening paradigm (LWL), which tracks eye movements to the referent of a word as a measure of online language comprehension. Yet the analytic and design choices for LWL vary widely, and the consequences of these choices have rarely been investigated at a large scale. Here we use data from Peekbank, a repository of LWL data from studies with almost 4000 children (6 months–5 years), to search for analytic and design decisions that maximize the reliability and validity of LWL measures. Analytic recommendations include using long analytic windows for computing accuracy and avoiding baseline correction. We also find little or no benefit from excluding children or trials with low data. The recommended combination of analytic choices increased reliability (intra-class correlations and test-retest correlations) by approximately .20 and validity (correlations with productive vocabulary) by .07. Effects were generally consistent across age, but gains for reliability were more substantial than for validity. Design recommendations include building experiments with a minimum of 10 target trials per child per condition and matching stimuli for animacy. We hope that these recommendations lead to improved and unified standards for the measurement of children’s word recognition in the LWL paradigm.

Keywords: looking while listening, word recognition, language development, eyetracking, reliability, validity

Word count: 7850

Data-driven recommendations for increasing reliability and validity in measures of infant word recognition

Introduction

Over the first few months of life, children become attuned to their native language, organizing the phonological system of speech input and parsing speech into meaningful units. Soon after, children begin to show early signs of comprehension of words or phrases through simple behaviors such as reaching for an object that is named (e.g., “Where’s the ball?”). These everyday accomplishments mark important milestones in children’s development. Yet global assessments of these milestones do not capture the millisecond-level processes that take place during real-time language comprehension. Children, like adults, process speech incrementally, as it unfolds in time (Fernald, Pinto, Swingley, Weinberg, & McRoberts, 1998; Frank et al., 2026; Snedeker, 2013), and use features of the speech signal to extract meaning that supports language comprehension (Law & Edwards, 2015; Weisleder & Fernald, 2013). One powerful, child-friendly method for tracking these moment-to-moment processes is precisely measuring looking behavior in response to speech. In particular, the “looking while listening” (LWL) procedure has emerged as a *de facto* standard for assessing language comprehension in infants and children (Frank, Braginsky, Yurovsky, & Marchman, 2021).¹ In the standard version of the LWL task, children see two pictures on a screen and hear speech directing their attention to one of the pictures, e.g., “look at the ball”. The accuracy and speed of fixating the named referent measure the efficiency of language comprehension (Fernald et al., 2008).

The LWL procedure has been a key tool for enhancing our understanding of the continuous processes of growth in incremental language processing (Golinkoff, Ma, Song, & Hirsh-Pasek, 2013). For example, using LWL methods, we can measure developmental increases in the speed and accuracy of children’s recognition of familiar words (Fernald, Perfors, & Marchman, 2006; Frank et al., 2026) and in their sensitivity to detailed phonological information during the first years after birth (Moore & Bergelson, 2024; Swingley & Aslin, 2000). Moreover, measures of children’s skill in the LWL task have proven useful as an index of individual differences. Children who show more efficient language processing of familiar words at 18 months show larger vocabulary size and more rapid vocabulary growth over the second and third years of life than children with less efficient processing (Fernald & Marchman, 2012; Fernald et al., 2006; Frank et al., 2026; Peter et al., 2019). Variation in language processing efficiency has also been shown to be associated with later language, cognitive, and non-verbal outcomes in both typically-developing (Law & Edwards, 2015; Peter et al., 2019) and preterm children (Loi, Marchman, Fernald, & Feldman, 2017; Marchman, Adams, Loi, Fernald, & Feldman, 2016; Marchman et al., 2018).

Despite widespread use of the LWL method in language development research, there is still limited consensus on how to best process and analyze time-series data of this type (from either adults or infants, Mirman, Dixon, & Magnuson, 2008; Bergelson & Swingley, 2012; Peelle & Van Engen, 2021; Von Holzen & Mani, 2012; Weaver, Saffran, & Arunachalam,

¹ Here we adopt LWL as a convenient label while recognizing that the procedure goes by a number of different names in the literature, including the “intermodal preferential looking procedure” (Fernald, Zangl, Portillo, & Marchman, 2008; Hirsh-Pasek & Golinkoff, 1996; Reznick, 1990).

2026). This lack of consensus likely stems from structural factors: infant recruitment is time- and cost-intensive and datasets collected by individual labs are often small (Bergmann et al., 2018; Frank, Braginsky, Yurovsky, & Marchman, 2017). These issues make it difficult for any one researcher to systematically evaluate the consequences of specific analytic choices independently of the data collected for a particular study (cf. Peelle & Van Engen, 2021). Instead, researchers make ad-hoc decisions and adopt idiosyncratic lab practices that themselves originate from subjective judgments of what “seems to work”. By decreasing uniformity in the literature, this variation also increases the challenge of triangulating what precisely is being measured by the LWL task (Weaver et al., 2026). How can we put the measurement and analysis of infant word recognition and online language processing on firmer empirical footing?

Our approach

Peekbank is a new data resource that offers a potential solution to this core issue by providing a large database of infant eye-tracking data from the LWL paradigm for exploring methodological questions (Zettersten et al., 2023). The most recent release of Peekbank (version 2026.1) houses 44 datasets from almost 4000 infants across a wide age range (6 months–5 years) and across a broad set of test words and materials. The goal of the current manuscript is to use these data to systematically explore key analytic decisions that researchers make when using this paradigm, offering best-practices guidance for measurement.

Our approach throughout is to seek analytic choices that distill the best possible measures from the rich time course of behavior in the LWL task. Recent reviews of the LWL literature highlight open questions about both the reliability and validity of looking-based measures of word recognition (Weaver et al., 2026). Starting with reliability, we aim for analyses that yield the maximal signal about individual children relative to both measurement error and irrelevant artifacts (Byers-Heinlein, Bergmann, & Savalei, 2022; Rhemtulla & Bottesini, 2022; Schreiner et al., 2024). Here, to measure reliability, we compute both within-dataset inter-item reliability and test-retest reliability for datasets with multiple LWL sessions close in time for the same children.

Turning to validity, the goal is to make analytic decisions that maximize the relationship with measures of word knowledge (Flake & Fried, 2020; Kominsky, Lucca, Thomas, Frank, & Hamlin, 2022). Here, we compute the correlations between LWL measures and vocabulary using aggregate measures of receptive and expressive vocabulary size from the MacArthur-Bates Communicative Development Inventories (MB-CDI). The MB-CDI is a popular parent report measure of early vocabulary (Frank et al., 2021; Marchman, Dale, & Fenson, 2023) that is widely used in conjunction with LWL tasks (Fernald & Marchman, 2012; Fernald et al., 2006; Marchman & Fernald, 2008). Although ideally we would have access to measures of precisely the same construct (e.g., button-press reaction times for word recognition), in practice, such data are not available at the scale we require and we thus rely on the CDI as a frequently-used measure of related but distinct constructs. Our focus in this paper is to provide guidance for navigating some of the many decisions facing an investigator analyzing LWL data. Previous work has provided in-depth recommendations for experiment design, eye-tracking, and eye-gaze coding (Fernald et al., 2008; Olson et al., 2020; Venker et

al., 2020; Wass, Smith, & Johnson, 2013). Other recent work using an earlier version of Peekbank (as well as other datasets) has also investigated the psychometric properties of a range of different dependent measures in LWL tasks and their general construct structure (Sander-Montant, Fibla, Killam, & Byers-Heinlein, 2026). We focus here instead on the consequences of varying analytic choices for two main measures which have the strongest psychometric properties (Sander-Montant et al., 2026). The first is *accuracy*: proportion looking to the target image (out of all looks to either the target or distractor image) computed over some period of time after the onset of the target word. The second measure is shift latency or *reaction time* (RT): how long it takes a child to shift their gaze from the distractor to the target image in response to the target word (Fernald et al., 2008).

Considering these measures, we address the following data processing (1-4) and design (5-6) questions with the goal of formulating broad recommendations. For each question, we also examine whether the results are modulated by age.

1. For accuracy, what start and end time maximize the signal of children’s response to the stimulus?
2. For accuracy, should analysts correct for baseline preferences in looking to the target image prior to the onset of the target word?
3. When computing RTs, is it best to operationalize RT as the latency from target word onset to when the child initiates an eye movement away from the distractor picture (launch RT) or to when a child completes this shift and fixates the target picture (landing RT)? What analytic window should be used? And should RTs be subject to a logarithmic transformation?
4. For both accuracy and RT, when should participants and trials be excluded from the analysis?
5. How many trials are needed to maximize reliability?
6. How do design choices in pairing target and distractor items affect measurement?

To foreshadow our results, we find that while many common data analysis practices provide reasonable levels of reliability and validity, it may be possible to further optimize reliability and validity through careful, data-driven decision-making. Overall, we find data-driven analytic choices that lead to gains in both reliability ($\sim .20$) and validity ($.07$); we return in the discussion to why gains in reliability might be more substantial.

Methods: Dataset and approach

Dataset

Peekbank is an aggregation of datasets from a number of LWL studies. This work is based on Peekbank version 2026.1 (released February 2026) which houses 44 datasets containing 134,929 trials from 3,915 children participating in 5,596 distinct testing sessions (some children participated in multiple sessions).

We limit our analyses to what we term “standard” trials, resulting in 75,636 trials from 3,171 children participating in 4,851 sessions across 36 datasets. For a trial to be considered “standard”, both the target and the distractor images must be of familiar objects. For the auditory stimulus, the target word must be a real word that occurs once and is the first and only source of information about the target. In our analyses, the target word is also always a noun. The carrier phrase must be a well-formed, grammatical phrase in a single language

(e.g., English, Spanish), without any intentionally-added noise or audio filtering. We focus on these trials because they are commonly used and most directly reflect the construct of incremental word recognition processes. (See SI: Criteria for Inclusion for more detailed criteria used to handle borderline situations.)

Following Peekbank’s convention, trial time is centered at the onset of the target word (labeled as 0 ms). We only include trials from children 60 months or younger; the average age in the final dataset was 22 months (SD: 9 months, median: 19 months; see Figure S1 for age histograms). The final dataset included 1,488 girls, 1,570 boys, and 113 children of unspecified gender. The majority of children were learning English (2,617 children), but some children were learning Spanish (154 children), English and French (200 children), Norwegian (69 children), Dutch (54 children), English and Spanish (54 children), and Tselal (23 children). No additional demographic information is available in the Peekbank database; more detailed demographic information on each dataset can be found in the corresponding papers.

MB-CDI comprehension scores were available for 632 children from 12 datasets, and MB-CDI production scores were available for 1,143 children from 20 datasets. Test-retest data come from 5 datasets with multiple sessions within a month of each other, covering 365 children (788 session pairs). See Table 1 for a list of included datasets. Figure S2 has time course plots showing the number of trials for each dataset.

Approach

Reliability refers to how consistently measurements of the same thing produce the same outcome. One measure of reliability is an intra-class correlation (ICC) that measures (within a dataset) how much different items (target labels; “rates” in ICC terminology) are consistent with each other in measuring a child within a session. Specifically, we computed ICCs based on a mean-rating, consistency, two-way random effects model (Type 2A in Gwet (2014)’s typology) using the agreement package in R (Girard, 2019) (see SI: Details about ICC for more specifics). Another measure of reliability is test-retest reliability, which measures how similarly the same child performs when assessed on two different occasions on the same task. This measure is only possible on a subset of the datasets that repeated sessions at closely spaced intervals. While test-retest reliability data are limited, this is the metric that is most important for measuring the stability of individual differences.

Validity refers to whether measures correspond to the construct of interest (in this case, language processing efficiency and/or receptive/expressive vocabulary size). We focus specifically on the relationships between aggregate measures of receptive and expressive vocabulary size from the MB-CDI and LWL task, given the documented relationships between these measures (e.g., Fernald & Marchman, 2012; Peter et al., 2019). This correspondence reflects the general relationship between familiar word recognition in the moment (indexed by LWL) and either concurrent vocabulary size or the growth of early vocabulary over time (indexed by the MB-CDI) (Frank et al., 2026).² Since some studies

² We do not explore relations at the item-level (e.g., speed of processing of a particular word and a child’s reported knowledge of that word) given the mixed evidence for item-level alignment between LWL and MB-CDI data (López Pérez, Moore, Sander-Montant, & Byers-Heinlein, 2024; Weaver & Saffran, 2026).

Table 1

Descriptives of the included datasets. Age is mean age in months, with the range shown in brackets.

	Dataset Name	Pct Trials	Trials	Sessions	Subjects	Age	CDI	Test-retest	Language
1	Adams et al. (2018)	18.4%	13,899	711	69	22 [13 - 38]	x	x	English
2	Fernald & Marchman (2012)	15.6%	11,828	679	122	22 [17 - 32]	x	x	English
3	Weaver et al. (2024)	6.5%	4,948	248	141	16 [14 - 24]		x	English
4	Fernald et al. (2013)	5.8%	4,359	178	80	20 [17 - 26]	x	x	English
5	Hurtado et al. (2008)	5.3%	4,035	136	76	21 [17 - 27]	x		Spanish
6	Fernald et al. (2006)	4.9%	3,692	229	63	20 [15 - 25]	x		English
7	Yurovsky et al. (2013)	3.7%	2,810	413	413	34 [12 - 60]			English
8	Yurovsky et al. (2024)	3.1%	2,350	327	327	37 [8 - 60]			English
9	Bergelson & Swingley (2012)	2.8%	2,081	84	84	12 [6 - 21]	x		English
10	Kartushina & Mayor (2019)	2.7%	2,072	69	69	7 [6 - 9]	x		Norwegian
11	Weaver & Saffran (2026)	2.4%	1,788	69	69	19 [18 - 20]	x		English
12	Weisleder & Fernald (2013)	2.3%	1,769	58	29	22 [18 - 27]			Spanish
13	Yurovsky & Frank (2017)	2.3%	1,747	285	285	26 [13 - 59]			English
14	Borovsky & Peters (2019)	2.2%	1,676	79	79	18 [17 - 20]	x		English
15	Potter & Lew-Williams (2024)	2.1%	1,626	67	67	24 [21 - 27]	x		English
16	Yoon et al. (2015)	2.0%	1,500	199	199	43 [13 - 60]			English
17	Sander-Montant et al. (2023)	1.8%	1,381	123	123	22 [12 - 31]	x		English, French
18	Mahr et al. (2015)	1.6%	1,232	29	29	21 [18 - 24]	x		English
19	Egger et al. (2020)	1.4%	1,052	54	54	18 [18 - 19]	x		Dutch
20	Garrison et al. (2020)	1.4%	1,046	35	35	15 [12 - 18]	x		English
21	Hurtado et al. (2007)	1.3%	1,010	49	49	24 [15 - 37]			Spanish
22	Pomper & Saffran (2017)	1.0%	785	82	82	17 [14 - 19]			English
23	Frank et al. (2016)	1.0%	781	108	108	35 [12 - 60]			English
24	Bacon & Saffran (2022)	1.0%	754	38	38	23 [22 - 24]	x		English
25	Ronfard et al. (2022)	1.0%	737	40	40	20 [18 - 24]	x		English
26	Perry & Saffran (2017)	0.9%	718	42	42	20 [19 - 22]	x		English
27	Casillas et al. (2017)	0.9%	651	23	23	32 [9 - 48]			Tseltal
28	Swingley & Aslin (2002)	0.8%	596	50	50	15 [14 - 16]	x		English
29	Pomper & Saffran (2016)	0.6%	471	60	60	44 [41 - 47]			English
30	Kremin et al. (2023)	0.5%	410	60	60	42 [37 - 49]			English, French, Spanish
31	Pomper & Saffran (2015)	0.5%	389	25	25	40 [38 - 42]			English
32	Moore & Bergelson (2022)	0.5%	384	29	29	18 [16 - 20]	x		English
33	Potter et al. (2019)	0.4%	309	44	23	23 [18 - 29]	x	x	English
34	Byers-Heimlein et al. (2017)	0.4%	284	48	48	20 [19 - 21]	x		English, French
35	Pomper & Saffran (2019)	0.3%	249	44	44	40 [38 - 43]			English
36	Pomper & Saffran (2018)	0.3%	217	37	37	39 [38 - 43]			English
	Total	100%	75,636	4,851	3,171	22 [6-60]	21	5	

used different versions of the CDIs (e.g., short vs. long forms), all scores were converted to a proportion of the possible items.

All code and data to reproduce the analyses are available at <https://github.com/peekbank/peekbank-method>.

Data processing recommendations

Recommendation 1: Measure accuracy over a long analytic window.

Accuracy, or proportion of target looking, must be computed over a time window of interest. Typically, windows start shortly after the onset of the target word and continue until a few seconds later; this window is intended to capture the peak of children’s stimulus-driven looks to the target. However, in practice, studies vary in what windows are used, and often window selection is a locus for analytic flexibility (choosing windows post hoc to maximize a particular effect) (Pelle & Van Engen, 2021). Common analysis windows include 300–1800 ms after target word onset (following Fernald et al., 2008) or 367–2000 ms after target word onset (e.g., Swingley & Aslin, 2002), although windows that extend until 3 or 4 seconds after target word onset are also used, especially at younger ages (e.g., Bergelson & Swingley, 2012; Bergelson & Aslin, 2017; López Pérez et al., 2024). In this analysis, we

begin from first principles and ask when windows should start and end to maximize reliability and validity on word recognition trials.




Figure 1. A heat map showing the average ICC across datasets for different start and end points for the accuracy window. A region with the highest ICC (best reliability) is outlined in black, and the region most commonly represented in the literature is outlined in blue.

Analytic approach. We first calculated the ICC reliability for accuracy across datasets, computing average ICC across a range of different analysis window start and end points. The standardized Peekbank format allowed us to apply each combination of window start and end points consistently within each dataset. Next, we added the additional metrics of test-retest reliability and – to assess validity – CDI correlations, which were calculated by dataset for each combination of parameter values. To aggregate across datasets, we used a linear mixed-effects model with random intercepts by dataset, weighted by the number of sessions. We interpret the fixed-effect contrasts as the predicted difference in outcome when switching from one set of parameters to another on a generic dataset. The best data-processing options might be different for children of different ages, so we ran the same linear mixed-effects models separately for children in each of four age groups: <18 months, 18–23 months, 24–35 months, and ≥ 36 months.

Results.

Reliability. Average ICCs across datasets for different windowing options are shown in Figure 1. As expected, including time points from before the target word onset decreases reliability. Our initial exploration found a region of high reliability with start times around 600–800 ms and end times around 2800–4000 ms. A dataset-by-dataset comparison of the differences between a 300–1800 ms window (representative of common practice) and a 600–4000 ms window are shown in Figure 2A. Nearly all datasets have higher ICC reliability for the later and longer window. All 5 test-retest reliability datasets also show increased reliability on this later and longer window. Our mixed-effects model confirms that there are reliability gains from switching from a 300–1800 ms window to a 600–4000 ms one. ICC reliability increased by 0.12 [0.08, 0.16] (square brackets indicate 95% CIs) and test-retest




Figure 2. Per-dataset outcome measures for two window options: 300–1800 ms (red) and 600–4000 ms (blue). Line color indicates whether the longer (blue) or shorter (red) window leads to a higher value. Dot size is proportional within each panel to the number of sessions.




Figure 3. Per-dataset ICC measures for two window options: 300–1800 ms (red) and 600–4000 ms (blue), split by age category. Line color indicates whether the longer (blue) or shorter (red) window leads to a higher value. Dot size is proportional within each panel to the number of sessions.

reliability by 0.07 [-0.02, 0.16].

These patterns generally hold across age. The overall level of reliability rises with age, as older children generally fixate target images more consistently and experiments with younger children typically use fewer trials (Figure S3). However, the overall pattern of which time windows maximize reliability at any given age is remarkably consistent across the different age groups (Figure S5). For instance, across different ages there is an improvement

in ICC from using a 600–4000 ms window compared to a 300–1800 ms one (Figure 3). For <18 months, the ICC difference is 0.06 [-0.03, 0.16]; for 18–23 months, 0.13 [0.07, 0.19]; for 24–35 months, 0.06 [0.01, 0.12]; and for ≥ 36 months 0.15 [0.06, 0.23].

Validity. There is no evidence for a difference in validity between a 300–1800 ms window and a 600–4000 ms window (Figure 2B). Switching to the longer window is associated with a change of -0.02 [-0.07, 0.03] in correlation with CDI comprehension and 0.03 [-0.02, 0.08] with CDI production.

Discussion. Overall, we find that LWL accuracy measures are reliable across a wide range of window choices; however, to maximize reliability, we recommend windows that start somewhere in the 400–600 ms range and extend to 3000–4000 ms (longer than the canonical window that ends at 1800 or 2000 ms), regardless of child age. Windowing choices seem to primarily affect reliability measures and have limited effects on validity measures. The finding that later starting and later ending windows increase reliability suggests that, even as looks to the target decline after 2000 ms, looking behavior still contains reliable signal. While some children may look quickly to the target and then look away, other children who look to the target later or who continue looking at the target are providing a useful signal during later portions of a given trial. Overall, children in many datasets continue to show consistent looking to the target over the distractor during later trial periods (Figure S2). An implication of these results is that choosing analysis windows based on the peak of proportion target looking in the averaged time course plot generally leads to lower reliability estimates, with no gains in validity.

Recommendation 2: Do not use baseline-corrected accuracy.

One commonly acknowledged issue in LWL studies is that looking behavior could be driven by differences in visual salience between the target and distractor images themselves (Fernald et al., 2008; Novack et al., 2025): Children may have a baseline preference to look at one image over another, even in the absence of auditory stimuli. Although thoughtful design choices can minimize these issues (see Recommendation 6), some analytic strategies also attempt to account for baseline preferences in dependent measures. While there are a variety of such strategies in the LWL literature, here we focus on one common practice: baseline correction. In baseline correction, researchers use a difference measure between accuracy on a critical post-stimulus window and “accuracy” (proportion looks to target among looks to target or distractor) on a pre-stimulus baseline window (e.g., Bergelson & Swingley, 2012; Weaver, Zettersten, & Saffran, 2024). The baseline window can be the 2 seconds prior to target word onset (Weaver et al., 2024) or the full 3–4 second pre-stimulus period (Bergelson & Swingley, 2012).

Analytic approach. We considered baseline-corrected accuracies with various pre-stimulus baseline windows, focusing on a 600–4000 ms critical window (although the results generalize to other critical windows). We restricted analysis to datasets with at least some data from 2000 ms or more before target word onset. We also computed a “no baseline” accuracy, based solely on accuracy within a 600–4000 ms time window with no correction. Baseline windows started at -1000, -2000, -3000 and -4000 ms and ended at target word onset.

Results.




Figure 4. Comparison between a baseline-corrected accuracy in red and accuracy in blue across datasets and measures. Both use a critical window of 600–4000 ms; the baseline window was -2000–0 ms. Line color indicates whether the longer (blue) or shorter (red) window leads to a higher value. Dot size is proportional within each panel to the number of sessions.

Reliability. Baseline-corrected accuracy led to worse reliability than critical-window accuracy alone. As an illustrative example, Figure 4A shows per-dataset results for accuracy compared to a baseline-corrected accuracy using a -2000–0 ms baseline window. This result held generally: All baseline window sizes that we considered led to worse reliability compared to using normal accuracy without any correction. For a critical window of 600–4000 ms, baseline correction decreased ICC by at least 0.10 [0.04, 0.16] and test-retest reliability by at least 0.19 [0.08, 0.30]. Similar patterns hold for different accuracy windows (Figure S6) and across different age groups (Figure S8).

Validity. Baseline-correction also reduces the validity of accuracy measures (Figure 4B). All baseline window sizes that we considered led to worse correlations with productive vocabulary (though results were less robust for receptive vocabulary). For a critical window of 600–4000 ms, baseline correction decreased the correlation with receptive vocabulary by at least 0.04 [-0.05, 0.12] and correlation with productive vocabulary by at least 0.09 [0.05, 0.13]. Similar patterns hold for different accuracy windows (Figure S6) and across different age groups (Figure S8).

Discussion. We do not recommend the use of baseline-correction in accuracy measures, as it worsens both reliability and validity compared to a critical-window-only accuracy measure. Why would incorporating baseline information lead to worse metrics? The literature on the psychometrics of difference scores (e.g., Lord, 1956; Rogosa & Willett, 1983; Trafimow, 2015) provides a valuable guide for evaluating the expected measurement properties of baseline correction. This literature points to two important features of LWL tasks: first, children’s looking behavior is noisy from trial to trial, meaning that the reliability of proportion looking during both the baseline and the critical window is low; and

second, the signal across baseline and critical portions of a trial is typically correlated (in part due to the very salience that the baseline correction is trying to remove). Unfortunately, it is under precisely these conditions – low reliability and high correlation between components – that difference scores exhibit the lowest reliability (Trafimow, 2015). Thus, while baseline-correction may lead to grand average accuracy scores that are less biased (by removing item salience), it also leads to less reliable and valid individual-level measures.

Recommendation 3: Measure landing reaction times with a minimum reaction time of around 400 ms.

Reaction time (RT) reflects the speed with which infants shift their gaze from the distractor to the target image. RT computations include many choice points. RTs can be computed based on when a child’s eyes leave the distractor (launch RT) or when a child’s eyes arrive on the target image (landing RT) (Fernald et al., 2008). A minimum and maximum RT may be set, as short and long RTs are often excluded because they are considered unlikely to be generated in response to the stimulus (Fernald et al., 2008). Reaction time data are sometimes analyzed as raw scores and sometimes analyzed using a logarithmic transform (consistent with the idea that reaction times are log-normally distributed and that effects on them may be multiplicative rather than additive) (Csibra, Hernik, Mascaro, Tatone, & Lengyel, 2016; Frank et al., 2026). We systematically investigated the effect of each of these choice points on reliability and validity.




Figure 5. Heat map of different parameter options for measuring RT and how they on average affected ICC. A region with the highest ICC (best reliability) is outlined in black, and the region most commonly represented in the literature is outlined in blue.

Analytic approach. We considered a wide range of options for measuring RT, fully crossing a) launch or landing RTs, b) a minimum RT threshold for inclusion between 0 and 1000 ms, c) a maximum RT threshold for inclusion between 1600 and 4000 ms, and d) raw or log-transformed RTs.

Results.




Figure 6. Comparison between two different sets of RT parameters. In red, an example of common practice, using a launch RT from 300–1800 ms. In blue, an RT that maximizes ICC, with a landing RT from 400–4000 ms, on the log scale. Note that correlations with the CDI are generally expected to be negative (as faster word recognition is associated with larger vocabulary). Line color indicates whether the longer (blue) or shorter (red) window leads to improved reliability or validity.

Reliability. An initial exploration indicated that ICC reliability was maximized for landing RTs with a minimum RT around 400 ms and a high maximum RT (Figure 5). A comparison between a common set of parameters for calculating RT and an alternative that maximizes ICC are shown in Figure 6A. Compared to computing the “standard” raw launch RTs restricted to 300–1800 ms, using log-scale landing RTs ranging between 400–4000 ms increased ICC by 0.11 [−0.01, 0.23], with little change in test-retest reliability (0.04 [−0.18, 0.27]). Similar patterns hold for children of different ages (Figure S9 and Figure S11).

Validity. Different methods for calculating RT had little impact on validity (Figure 6B). Compared to computing the “standard” raw launch RTs restricted to 300–1800 ms, using landing log RTs ranging between 400–4000 ms has little impact on CDI correlations (−0.04 [−0.17, 0.08] for comprehension, and −0.03 [−0.12, 0.06] for production).

Discussion. Overall, RT measures capture substantial signal, as indicated by the moderate ICC values across all choice parameters. A wide range of parameter options have similar properties (Figure S10 and Figure S12), but computing landing RTs with an inclusion window of 400–4000 ms maximizes reliability without negatively impacting validity. Note that the recommendation to compute landing reaction times contradicts traditional approaches to reaction-time computation (Fernald et al., 2008) based on launch time; however, landing RTs generally lead to higher ICCs in our dataset. Log-scale transformations should be considered, but may depend on modeling and theoretical considerations.

Recommendation 4: Trial-level and session-level exclusions for low amounts of data are generally unnecessary.

A common issue in developmental research is when to apply trial- and session-level exclusions. For LWL studies, trial-level exclusions tend to be applied when children contribute too few time points (e.g., due to tracker loss or looking away from the screen during the trial). Papers often use a combination of heuristic exclusions for inattention prior to coding looking behavior (Adams et al., 2018) and criterion-based exclusions based on the proportion of looking to the target and distractor during the trial. Cutoffs vary, with different experiments requiring minimum on-screen looking for at least 20%, 50%, or 66% of the critical window of a trial (Borovsky & Peters, 2019; Pomper & Saffran, 2019; Weaver et al., 2024). Studies also occasionally filter for on-screen looks specifically at the point of disambiguation (C. E. Potter, Fourakis, Morin-Lessard, Byers-Heinlein, & Lew-Williams, 2019).

In addition, studies also implement session-level exclusions based on the minimum number of valid trials. For accuracy measures, children are often excluded if they had fewer than 25% or 50% of the total trials in a session, either within a condition or overall (Adams et al., 2018; Garrison, Baudet, Breitfeld, Aberman, & Bergelson, 2020; Perry & Saffran, 2017). For RT, it is instead standard to require a minimum number of trials (e.g., two, Fernald et al., 2006; Adams et al., 2018).

There are multiple possible ways that trial- and session-level exclusions could affect measurement properties. If trials with fewer valid looking points and sessions with only a small number of trials still provide signal of the child-level language behavior, excluding them will decrease reliability and validity. On the other hand, if these trials and sessions reflect noise, excluding them will increase reliability and validity.

Analytic approach. We consider 3 types of exclusions: trial-level exclusions for accuracy, session-level exclusions for accuracy, and session-level exclusions for RT. We first explored trial-level accuracy exclusions, both on the basis of proportion of off-screen looks and on the basis of where the child was looking at target word onset. For accuracy, we then considered per-session minimums either based on the number of valid trials, or based on the fraction of valid trials in a session. The fraction of valid trials is the more common exclusion in the literature, but when drawing from many datasets that vary in the total number of trials, minimum trial count likely represents a more comparable metric. For RT, we considered different cutoffs for the minimum number of valid trials required for inclusion.

We note that datasets contributed to Peekbank vary in whether excluded data was contributed, or whether only cleaned post-exclusion data were contributed. Thus, some datasets may not show effects of exclusions if those trials or sessions that would be excluded are missing from Peekbank.

Results.

Reliability. For trial-level exclusions, a requirement to look on-screen at least 50% of the time results in 85.64% (range: 27.22%–98.92%) of trials being retained. These exclusions increased in ICC by 0.03 [0.00, 0.05] and no substantial change in test-retest reliability (-0.02 [-0.11, 0.08]). Filtering based on initial looking position, or having requirements for looking stricter than 50%, start to diminish the amount of data and potentially negatively affect test-retest reliability (Figure S13 and Figure S14).

We then explored the effects of excluding sessions for too few trials. Here, we focus on the results for a minimum number of 8 total trials, with no trial-level exclusions, to illustrate general patterns in the results, but report outcomes for a range of thresholds in the supplement (Figures S15–S18). Only including sessions with at least 8 trials results in retaining 80.99% (range: 0%–100%) of total sessions³. These exclusions show little evidence of increasing ICC (0.00 [-0.01, 0.02]) or test-retest reliability (0.03 [-0.08, 0.14]) compared to retaining all sessions.

We also considered the effects of minimum trial requirements on RT data. Only including sessions that contribute at least 2 trials results in retaining 90.93% (range: 42.86%–100%) of all sessions. Similar to previous results, these exclusions lead to small changes in reliability (ICC: -0.01 [-0.07, 0.04]; test-retest reliability: 0.03 [-0.08, 0.14]).

Validity. Excluding trials with less than 50% on-screen looking did not substantially change the validity of accuracy measures compared to no exclusions (comprehension: 0.01 [-0.04, 0.06]; production: 0.00 [-0.03, 0.03]). As with reliability, on-screen looking criteria stricter than 50% increase data loss and consequently may even have negative effects on validity (Figure S13 and Figure S14). Similarly, session-level exclusions did not substantially impact validity. Only including sessions with at least 8 trials, compared to including all sessions, led to virtually no increase in correlations (comprehension: 0.01 [-0.04, 0.07]; production: 0.01 [-0.03, 0.05]). For RT, excluding sessions with fewer than 2 valid RT trials also led to no substantive changes in correlations with CDI (comprehension: 0.00 [-0.15, 0.14]; production: -0.02 [-0.08, 0.04]).

Discussion. Standard practice within the field often errs on the side of excluding a large number of trials and sessions for missing data. Surprisingly, our results show that applying no exclusions or minimal exclusions is often a reasonable choice, at least in terms of the psychometric properties of the resulting measurements.

Design recommendations

Our focus for this work was to explore how data processing choices affected the reliability and validity of the resultant measures. However, along the way, we identified several experimental design questions that had notable consequences for our central measures. We summarize two design recommendations that stood out as having the largest potential impact on reliability and validity (see supplement, section S7 for an additional analysis of order effects).

Recommendation 5: Design studies to have a minimum of 10 trials per child, per condition, or as many trials as possible for individual-difference studies.

We found limited evidence for excluding children for low trial counts, but intuitively, more repeated measures should improve measurement properties such as reliability and validity. Rather than retrospective exclusions, another approach is prospective: how many trials per child should one aim for to achieve adequate reliability and validity? Test-retest reliability is important for studies that investigate longitudinal effects, as it assumes that the

³ The number of standard trials in a session varied, depending on overall number of trials and whether non-standard trials were also present.




Figure 7. Simulations of measurement properties for different trial counts, created by downsampling trials for all children with different trial minimums. Shown for an accuracy window of 600–4000 ms with no trial- or administration- level exclusions.

test sessions are representative of properties of the child at that age (Byers-Heinlein et al., 2022). To address this prospective question, we considered how reliability and validity measures would have changed if we had fewer trials per child. We filtered the data to include only children who had contributed at least N trials (for different values of N) and then repeatedly subsampled their trials, calculating metrics for each subset. Datasets vary in how many trials they have, so different minimum trials include different distributions and datasets, items, and children.

The results for accuracy and RT are shown in Figure 7. ICC reliability is not sensitive to the number of trials per subject, so it is not shown. For accuracy, more trials per session increase test-retest reliability, up to 20 trials (the highest we measured), although the gains per additional trial decrease with more trials. CDI-based validity also improved as trials increased, with smaller gains past 5 trials. Based on accuracy, we recommend at least 10 trials per condition in general, and as many as possible for individual differences studies where test-retest reliability is essential.

Similar considerations apply for RT, but at lower total trial numbers. Nearly every trial a child does results in a usable accuracy measure, but only around 1/3 of trials result in a valid RT measure in our data. For test-retest reliability, having at least 5 RT measures per session leads to notable improvements; gains beyond 5 are hard to measure because of the low number of subjects with sufficient trials, but more trials appear to be beneficial. For CDI measures, there are substantial gains for having 2 trials compared to 1 and continued

benefits to including additional trials. Overall, for measuring individual differences in processing speed, designing studies to include as many trials as possible appears to be central to adequate measurement.

Recommendation 6: Match target and distractor images based on animacy.

One concern with LWL trial design is that children may have preferences for which image to look at separate from the auditory stimuli, and some of their looking behavior may be independent of the auditory stimulus. One approach is to decrease the preference discrepancy between a pair of images, so that in the absence of auditory cues, children prefer to look at both images approximately equally. While there are a number of factors that could play into children’s image preferences (Fernald et al., 2008; Nothdurft, 2000; Pomper & Saffran, 2019), one factor that stood out in the data was animacy. Children have a strong preference for looking at animals and humans. When trials have one animate and one inanimate, children look more at the animate in the pre-stimulus period (Figure 8A) and are slower to respond when the inanimate is named. In contrast, for trials with target and distractor items matched on animacy, children’s early looking patterns for animate and inanimate targets are highly similar. This difference is particularly relevant for the younger children, whose stimulus-contingent looking behavior is weakest, so more of the looking behavior is driven by animacy; however, effects of animacy were apparent in children up to the age of three (Figure 8B).

This finding has important consequences for estimating word recognition: accuracy for animate targets is substantially overestimated and accuracy for inanimate targets is substantially underestimated when animacy is not balanced in the target-distractor pairing. The results are especially dramatic for younger infants (under 18 months): looking to familiar inanimate targets can appear to be at chance when paired with an animate distractor. In other words, an improper balancing of saliency could make an infant appear not to know a word, even though they do.

How much can this matter?

To quantify the potential cumulative consequences of different combinations of choices, we computed the difference in reliability when making all of the recommended choices with respect to data processing to an instance in which a researcher made a set of unfortunate choices that are nevertheless consistent with the literature. Consider two researchers analyzing accuracy in the same LWL dataset, Norm and Rex. Researcher Norm makes a set of reasonable choices based on practices observed in other LWL papers, choosing to analyze a short analysis time window (300–1800 ms); apply baseline correction (baseline window: -2000–0 ms); exclude trials on which infants are not looking at the screen for at least 60% of the critical window or not looking at the screen during target word onset; and require that infants contribute at least 50% of all trials for inclusion. Researcher Rex, on the other hand, analyzes the same data and instead chooses parameters based on the data-driven recommendations above: a critical window of 600–4000 ms, no baseline correction, and including all trials and sessions. Researcher Rex would see substantial increases in ICC (0.20 [0.13, 0.27]) and test-retest reliability (0.21 [0.13, 0.29]). Additionally, Rex could expect small increases in validity for correlations with comprehension (0.07 [-0.03, 0.17]) and




Figure 8. Looking time profiles for different combinations of target and distractor animacy. Panel A shows overall; Panel B shows split by age category.




Figure 9. Differences in dependent measures between the recommended accuracy measurement and a sensible-sounding, but unfortunate alternative set of parameters.

production (0.07 [0.00, 0.13]). Despite both researchers making choices consistent with past practices, we would expect Rex to derive a measure with improved psychometric properties compared to Norm, and therefore be in a better position to study individual differences and longitudinal change.

General Discussion

Children’s early word recognition is a central construct for the study of language learning, so improving recognition measures can have widespread benefits for our

understanding of language development. How can we maximize reliable, valid signal in the analysis of data in the LWL paradigm? We took a data-driven approach to this question, using a large-scale dataset of infant LWL data to quantify the impact of different analysis choices on measurement properties, leading to six core recommendations for data processing and experimental design.

For data processing, our strongest take-homes concern window choice and baseline correction. First, longer time windows (e.g., windows that end around 3000–4000 ms after target word onset) maximize the reliability of accuracy measures, with no negative consequences for validity. While average accuracy tends to peak earlier in the timecourse of a trial (typically 1000–2000 ms after target word onset), these findings suggest that later looking windows capture signal that boosts the consistency of word recognition measures. Second, while baseline correction is an intuitive solution to imbalances in the saliency of images, it causes large decrements to reliability and moderate decrements to validity; we therefore recommend avoiding this practice. Finally, we formulate recommendations for choosing reaction time cutoffs (minimum cutoffs around 400 ms tend to maximize reliability) and find that implementing trial- and session-level exclusions based on the amount of data contributed tend to have minimal consequences. Overall, choices about windows of analysis, approaches to exclusions, and how scores are computed can all have meaningful impacts on measure reliability.

One question that emerges from our analysis concerns the nature of the construct that LWL measures. In our analysis, long accuracy windows – as well as our other recommendations – increased reliability more than validity. Without more detailed characterizations of individual children, we cannot offer definitive answers about what drives this discrepancy, but there are several possible, non-mutually exclusive, explanations. One possibility is simply that some aspects of this difference relate to measurement error in CDI scores, which attenuates validity. Second, although the CDI is one useful source of validity evidence, it is possible that there are components of in-the-moment language processing that are not related to the primary construct captured by the CDI, which is general word knowledge. Finally, perhaps some aspects of our recommendations boost reliability in part because they capture non-linguistic variation between children. For example, perhaps signal from later in the accuracy timecourse (which boosts reliability) might be less coupled to language *per se*, and might instead reflect more general features of attention or processing speed. Future research could use this difference as a starting point to further parcellate the attentional and linguistic contributors to language-driven eye-movements.

A more general take-home from our data processing analyses is that looking at patterns in grand-average time courses is not a trustworthy guide to child- or item-level psychometric properties. Past recommendations for time window selection have sometimes been motivated by inspecting consistent patterns of looking in grand averages and noting when looking begins to decline across participants. These heuristics – while based on important patterns in looking behavior – may not always lead to optimal choices with respect to core measurement properties such as reliability. Wherever possible, measurement choices should be guided by direct estimates of psychometric properties in order to maximize the signal in infant looking behavior (Byers-Heinlein et al., 2022; Sander-Montant et al., 2026; Weaver et al., 2026).

Among our design-focused recommendations, the most striking result is the impact of animacy on infant looking behavior. When an image of an animate referent was paired with

an image of an inanimate referent, infant looking was systematically biased toward the animate image, even after target word onset. These results emphasize the importance of balancing saliency in image pairings as much as possible in LWL designs, particularly given that many of the current analytic approaches to correcting for saliency imbalances post hoc rely on difference-based measures which decrease reliability.

The second design recommendation is to include as many trials as possible when estimating effects for individual children. These findings converge with other recent results in the developmental literature emphasizing the importance of increasing per-participant observations in individual differences research (Byers-Heinlein et al., 2022; DeBolt, Rhemtulla, & Oakes, 2020; Schreiner et al., 2024).

Limitations & Generalizability

The scope of our recommendations is fundamentally limited by the characteristics of the Peekbank dataset. First, while the Peekbank database is by far the largest dataset of its kind, it is not a comprehensive or representative set of LWL studies. In particular, (1) the data are predominantly English word recognition trials collected from U.S. infants, and (2) the dataset is a convenience sample from labs who made their data available to the Peekbank team. LWL studies as a whole are likely more heterogeneous than what is represented in Peekbank – the datasets considered here come predominantly from a small network of researchers with similar training in LWL data collection.

At the same time, there is also substantial heterogeneity in the studies included in Peekbank, as different datasets were designed with different research questions in mind. To mitigate this concern, our analyses focused exclusively on standard trials that present familiar words with familiar images with limited or no targeted manipulations of the stimuli. The main benefit of focusing on standard trials is that it allows for a clear interpretation of reliability and validity – trials are truly designed to measure word recognition, rather than some other aspect of language processing. We view the remaining heterogeneity of study properties (e.g., specific words; differences in age) as a strength of our approach – our findings generalize broadly across a wide range of items and ages. However, by focusing on standard trials only, we also limit our conclusions to measuring familiar word recognition; our findings do not address how data processing and design practices would generalize to other estimation goals, e.g., estimating the magnitude of differences between two experimental conditions. Some of the more consequential choices, e.g., baseline correction, are likely to have a similar impact given general analysis principles known to affect reliability (e.g., computing difference scores). In general, our recommendations can be seen as a reasonable default choice when analyzing LWL data in the absence of data-driven guidance tuned to one’s specific estimation goal.

There is no one choice of time window that we can definitively identify as the “best” for maximizing reliability – instead, we see a pattern of graded improvements in reliability that follow a predictable pattern. While we sometimes articulate which choices lead to the highest reliability descriptively, these are at best approximations that we offer as data-driven default choices.

Conclusions

Experimental design and data processing decisions can have profound impacts on the properties of infant measures. Our findings show how compiling large-scale datasets such as Peekbank can help us estimate the impact of numerous experimenter decisions, generalizable across a wide range of ages, stimuli, and experimental contexts. Establishing consistent practices and principles for reliable and valid measurement of infant looking sets the stage for asking central questions about how language skill develops.

References

- Adams, K. A., Marchman, V. A., Loi, E. C., Ashland, M. D., Fernald, A., & Feldman, H. M. (2018). Caregiver talk and medical risk as predictors of language outcomes in full term and preterm toddlers. *Child Development, 89*(5), 1674–1690.
- Bacon, D., & Saffran, J. (2022). Role of speaker gender in toddler lexical processing. *Infancy : The Official Journal of the International Society on Infant Studies, 27*(2), 291–300.
- Bergelson, E., & Aslin, R. (2017). Semantic Specificity in One-Year-Olds' Word Comprehension. *Language Learning and Development, 13*(4), 481–501.
<https://doi.org/10.1080/15475441.2017.1324308>
- Bergelson, E., & Swingley, D. (2012). At 6-9 months, human infants know the meanings of many common nouns. *Proceedings of the National Academy of Sciences, 109*(9), 3253–3258.
- Bergmann, C., Tsuji, S., Piccinini, P. E., Lewis, M. L., Braginsky, M., Frank, M. C., & Cristia, A. (2018). Promoting replicability in developmental research through meta-analyses: Insights from language acquisition research. *Child Development, 89*(6), 1996–2009.
- Borovsky, A., & Peters, R. E. (2019). Vocabulary size and structure affects real-time lexical recognition in 18-month-olds. *PLOS ONE, 14*(7), e0219290.
<https://doi.org/10.1371/journal.pone.0219290>
- Byers-Heinlein, K., Bergmann, C., & Savalei, V. (2022). Six solutions for more reliable infant research. *Infant and Child Development, 31*(5), e2296. <https://doi.org/10.1002/icd.2296>
- Byers-Heinlein, K., Morin-Lessard, E., & Lew-Williams, C. (2017). Bilingual infants control their languages as they listen. *Proceedings of the National Academy of Sciences, 114*(34), 9032–9037.
- Casillas, M., Brown, P., & Levinson, S. C. (2017). *Casillas homebank corpus*.
<https://doi.org/10.21415/T51X12>
- Csibra, G., Hernik, M., Mascaro, O., Tatone, D., & Lengyel, M. (2016). Statistical treatment of looking-time data. *Developmental Psychology, 52*(4), 521–536.
- DeBolt, M. C., Rhemtulla, M., & Oakes, L. M. (2020). Robust data and power in infant research: A case study of the effect of number of infants and number of trials in visual preference procedures. *Infancy, 25*(4), 393–419.
<https://doi.org/https://doi.org/10.1111/inf.12337>
- Egger, J., Rowland, C. F., & Bergmann, C. (2020). Improving the robustness of infant lexical processing speed measures. *Behavior Research Methods, 52*(5), 2188–2201.
- Fernald, A., & Marchman, V. A. (2012). Individual differences in lexical processing at 18 months predict vocabulary growth in typically developing and late-talking toddlers. *Child Development, 83*(1), 203–222.
- Fernald, A., Marchman, V. A., & Weisleder, A. (2013). SES differences in language processing skill and vocabulary are evident at 18 months. *Developmental Science, 16*(2), 234–248.
- Fernald, A., Perfors, A., & Marchman, V. A. (2006). Picking up speed in understanding: Speech processing efficiency and vocabulary growth across the 2nd year. *Developmental Psychology, 42*(1), 98–116. <https://doi.org/10.1037/0012-1649.42.1.98>
- Fernald, A., Pinto, J. P., Swingley, D., Weinberg, A., & McRoberts, G. W. (1998). Rapid

- gains in speed of verbal processing by infants in the 2nd year. *Psychological Science*, 9(3), 228–231.
- Fernald, A., Zangl, R., Portillo, A. L., & Marchman, V. A. (2008). Looking while listening: Using eye movements to monitor spoken language comprehension by infants and young children. In I. A. Sekerina, E. M. Fernández, & H. Clahsen (Eds.), *Language Acquisition and Language Disorders* (Vol. 44, pp. 97–135). Amsterdam: John Benjamins Publishing Company. <https://doi.org/10.1075/lald.44.06fer>
- Flake, J. K., & Fried, E. I. (2020). Measurement Schmeasurement: Questionable Measurement Practices and How to Avoid Them. *Advances in Methods and Practices in Psychological Science*, 3(4), 456–465. <https://doi.org/10.1177/2515245920952393>
- Frank, M. C., Braginsky, M., Yurovsky, D., & Marchman, V. A. (2017). Wordbank: An open repository for developmental vocabulary data. *Journal of Child Language*, 44(3), 677–694.
- Frank, M. C., Braginsky, M., Yurovsky, D., & Marchman, V. A. (2021). *Variability and consistency in early language learning: The wordbank project*. Cambridge, MA: MIT Press.
- Frank, M. C., Marchman, V. A., Bergey, C. A., Boyce, V., Braginsky, M., Kachergis, G., . . . Zettersten, M. (2026). Continuous developmental changes in word recognition support language learning across early childhood. *eLife*. <https://doi.org/10.7554/eLife.109636.2>
- Frank, M. C., Sugarman, E., Horowitz, A. C., Lewis, M. L., & Yurovsky, D. (2016). Using tablets to collect data from young children. *Journal of Cognition and Development*, 17(1), 1–17.
- Garrison, H., Baudet, G., Breitfeld, E., Aberman, A., & Bergelson, E. (2020). Familiarity plays a small role in noun comprehension at 12–18 months. *Infancy : The Official Journal of the International Society on Infant Studies*, 25(4), 458–477.
- Girard, J. (2019). *Agreement: R package for the tidy calculation of inter-rater reliability*. Retrieved from <https://github.com/jmgirard/agreement>
- Golinkoff, R. M., Ma, W., Song, L., & Hirsh-Pasek, K. (2013). Twenty-five years using the intermodal preferential looking paradigm to study language acquisition: What have we learned? *Perspectives on Psychological Science*, 8(3), 316–339.
- Gwet, K. L. (2014). *Handbook of inter-rater reliability: The definitive guide to measuring the extent of agreement among raters* (4th ed.). Gaithersburg, MD: Advanced Analytics, LLC.
- Hirsh-Pasek, K., & Golinkoff, R. M. (1996). The intermodal preferential looking paradigm: A window onto emerging language comprehension. In *Language, Speech, and Communication. Methods for assessing children's syntax* (pp. 105–124). Cambridge, MA, US: The MIT Press.
- Hurtado, N., Marchman, V. A., & Fernald, A. (2007). Spoken word recognition by Latino children learning Spanish as their first language. *Journal of Child Language*, 34(2), 227–249. <https://doi.org/10.1017/S0305000906007896>
- Hurtado, N., Marchman, V. A., & Fernald, A. (2008). Does input influence uptake? Links between maternal talk, processing speed and vocabulary size in spanish-learning children. *Developmental Science*, 11(6), F31–F39.
- Kartushina, N., & Mayor, J. (2019). Word knowledge in six-to nine-month-old norwegian infants? Not without additional frequency cues. *Royal Society Open Science*, 6(9),

180711.

- Kominsky, J. F., Lucca, K., Thomas, A. J., Frank, M. C., & Hamlin, J. K. (2022). Simplicity and validity in infant research. *Cognitive Development*, 63, 101213. <https://doi.org/10.1016/j.cogdev.2022.101213>
- Kremin, L. V., Jardak, A., Lew-Williams, C., & Byers-Heinlein, K. (2023). Bilingual children's comprehension of code-switching at an uninformative adjective. *Language Development Research*, 3(1), 249.
- Law, F., II, & Edwards, J. R. (2015). Effects of Vocabulary Size on Online Lexical Processing by Preschoolers. *Language Learning and Development*, 11(4), 331–355. <https://doi.org/10.1080/15475441.2014.961066>
- Loi, E. C., Marchman, V. A., Fernald, A., & Feldman, H. M. (2017). Using eye movements to assess language comprehension in toddlers born preterm and full term. *The Journal of Pediatrics*, 180, 124–129. <https://doi.org/10.1016/j.jpeds.2016.10.004>
- López Pérez, M., Moore, C., Sander-Montant, A., & Byers-Heinlein, K. (2024). Infants' Knowledge of Individual Words: Investigating Links Between Parent Report and Looking Time. *Infancy*, 30(1), e12641. <https://doi.org/10.1111/inf.12641>
- Lord, F. M. (1956). The Measurement of Growth. *Educational and Psychological Measurement*, 16(4), 421–437. <https://doi.org/10.1177/001316445601600401>
- Mahr, T., McMillan, B. T. M., Saffran, J. R., Ellis Weismer, S., & Edwards, J. (2015). Anticipatory coarticulation facilitates word recognition in toddlers. *Cognition*, 142, 345–350. <https://doi.org/10.1016/j.cognition.2015.05.009>
- Marchman, V. A., Adams, K. A., Loi, E. C., Fernald, A., & Feldman, H. M. (2016). Early language processing efficiency predicts later receptive vocabulary outcomes in children born preterm. *Child Neuropsychology*, 22(6), 649–665. <https://doi.org/10.1080/09297049.2015.1038987>
- Marchman, V. A., Dale, P., & Fenson, L. (2023). *MacArthur-Bates communicative development inventories user's guide and technical manual*. Brookes Publishing.
- Marchman, V. A., & Fernald, A. (2008). Speed of word recognition and vocabulary knowledge in infancy predict cognitive and language outcomes in later childhood. *Developmental Science*, 11(3), F9–F16.
- Marchman, V. A., Loi, E. C., Adams, K. A., Ashland, M., Fernald, A., & Feldman, H. M. (2018). Speed of language comprehension at 18 months old predicts school-relevant outcomes at 54 months old in children born preterm. *Journal of Developmental & Behavioral Pediatrics*, 39(3), 246–253.
- Mirman, D., Dixon, J. A., & Magnuson, J. S. (2008). Statistical and computational models of the visual world paradigm: Growth curves and individual differences. *Journal of Memory and Language*, 59(4), 475–494. <https://doi.org/10.1016/j.jml.2007.11.006>
- Moore, C., & Bergelson, E. (2022). Examining the roles of regularity and lexical class in 18–26-month-olds' representations of how words sound. *Journal of Memory and Language*, 126, 104337.
- Moore, C., & Bergelson, E. (2024). Wordform variability in infants' language environment and its effects on early word learning. *Cognition*, 245, 105694. <https://doi.org/10.1016/j.cognition.2023.105694>
- Nothdurft, H.-C. (2000). Salience from feature contrast: Additivity across dimensions. *Vision Research*, 40(10-12), 1183–1201.

- Novack, M. A., Dworak, E. M., Han, Y. C., Kaat, A. J., Ustsinovich, V., Saffran, J., . . . Gershon, R. C. (2025). Development and validation of the NIH Baby Toolbox® language measures. *Infant Behavior and Development*, 80, 102121. <https://doi.org/10.1016/j.infbeh.2025.102121>
- Olson, R. H., Pomper, R., Potter, C. E., Hay, J. F., Saffran, J. R., Ellis Weismer, S., & Lew-Williams, C. (2020). *Peyecoder: An open-source program for coding eye movements*. <https://doi.org/10.5281/zenodo.3939234>
- Peelle, J. E., & Van Engen, K. J. (2021). Time Stand Still: Effects of Temporal Window Selection on Eye Tracking Analysis. *Collabra: Psychology*, 7(1), 25961. <https://doi.org/10.1525/collabra.25961>
- Perry, L. K., & Saffran, J. R. (2017). Is a pink cow still a cow? Individual differences in toddlers' vocabulary knowledge and lexical representations. *Cognitive Science*, 41(4), 1090–1105. <https://doi.org/10.1111/cogs.12370>
- Peter, M. S., Durrant, S., Jessop, A., Bidgood, A., Pine, J. M., & Rowland, C. F. (2019). Does speed of processing or vocabulary size predict later language growth in toddlers? *Cognitive Psychology*, 115, 101238.
- Pomper, R., & Saffran, J. R. (2015). *Unpublished "prime" study: Modulating attention to different features of objects during word learning*.
- Pomper, R., & Saffran, J. R. (2016). Roses are red, socks are blue: Switching dimensions disrupts young children's language comprehension. *PloS One*, 11(6), e0158459. <https://doi.org/10.1371/journal.pone.0158459>
- Pomper, R., & Saffran, J. R. (2017). *Unpublished "dimy" study: Do infants learn to associate diminutive forms with animates?*
- Pomper, R., & Saffran, J. R. (2018). *More than distractors: Familiar objects influence toddlers' semantic representations in novel word learning*. [Poster]. Poster presented at the XXI biennial international congress of infant studies., Philadelphia.
- Pomper, R., & Saffran, J. R. (2019). Familiar object salience affects novel word learning. *Child Development*, 90(2), e246–e262. <https://doi.org/10.1111/cdev.13053>
- Potter, Christine E., & Casey, L.-W. (2024). Frequent vs. Infrequent words shape toddlers' real-time sentence comprehension. *Journal of Child Language*, 51(6), 1478–1488.
- Potter, C. E., Fourakis, E., Morin-Lessard, E., Byers-Heinlein, K., & Lew-Williams, C. (2019). Bilingual toddlers' comprehension of mixed sentences is asymmetrical across their two languages. *Developmental Science*, 22(4), e12794. <https://doi.org/10.1111/desc.12794>
- Reznick, J. S. (1990). Visual preference as a test of infant word comprehension. *Applied Psycholinguistics*, 11(2), 145–166.
- Rhemtulla, M., & Bottesini, J. G. (2022). On the reliability of reliability estimates in infancy research. *Infant and Child Development*, 31(5), e2327. <https://doi.org/10.1002/icd.2327>
- Rogosa, D. R., & Willett, J. B. (1983). Demonstrating the Reliability the Difference Score in the Measurement of Change. *Journal of Educational Measurement*, 20(4), 335–343. <https://doi.org/10.1111/j.1745-3984.1983.tb00211.x>
- Ronfard, S., Wei, R., & Rowe, M. L. (2022). Exploring the linguistic, cognitive, and social skills underlying lexical processing efficiency as measured by the looking-while-listening paradigm. *Journal of Child Language*, 49(2), 302–325. <https://doi.org/10.1017/S0305000921000106>

- Sander-Montant, A., Fibla, L., Killam, H., & Byers-Heinlein, K. (2026). Keeping an eye on looking measures: Towards more robust developmental methods. *Developmental Science*, 29(4), e70228.
- Sander-Montant, A., Pérez, M. L., & Byers-Heinlein, K. (2023). The more they hear the more they learn? Using data from bilinguals to test models of early lexical development. *Cognition*, 238, 105525. <https://doi.org/10.1016/j.cognition.2023.105525>
- Schreiner, M. S., Zettersten, M., Bergmann, C., Frank, M. C., Fritzsche, T., Gonzalez-Gomez, N., et al.others. (2024). Limited evidence of test-retest reliability in infant-directed speech preference in a large preregistered infant experiment. *Developmental Science*, 27(6), e13551. <https://doi.org/10.1111/desc.13551>
- Snedeker, J. (2013). Children's sentence processing. In *Sentence Processing*. Psychology Press.
- Swingle, D., & Aslin, R. N. (2000). Spoken word recognition and lexical representation in very young children. *Cognition*, 76(2), 147–166. [https://doi.org/10.1016/s0010-0277\(00\)00081-0](https://doi.org/10.1016/s0010-0277(00)00081-0)
- Swingle, D., & Aslin, R. N. (2002). Lexical Neighborhoods and the Word-Form Representations of 14-Month-Olds. *Psychological Science*, 13(5), 480–484. <https://doi.org/10.1111/1467-9280.00485>
- Trafimow, D. (2015). A defense against the alleged unreliability of difference scores. *Cogent Mathematics*, 2(1), 1064626. <https://doi.org/10.1080/23311835.2015.1064626>
- Venker, C. E., Pomper, R., Mahr, T., Edwards, J., Saffran, J., & Ellis Weismer, S. (2020). Comparing automatic eye tracking and manual gaze coding methods in young children with autism spectrum disorder. *Autism Research*, 13(2), 271–283. <https://doi.org/https://doi.org/10.1002/aur.2225>
- Von Holzen, K., & Mani, N. (2012). Language nonselective lexical access in bilingual toddlers. *Journal of Experimental Child Psychology*, 113(4), 569–586. <https://doi.org/10.1016/j.jecp.2012.08.001>
- Wass, S. V., Smith, T. J., & Johnson, M. H. (2013). Parsing eye-tracking data of variable quality to provide accurate fixation duration estimates in infants and adults. *Behavior Research Methods*, 45, 229–250.
- Weaver, H., & Saffran, J. R. (2026). Interrogating Early Word Knowledge: Factors That Influence the Alignment Between Caregiver-Report and Experimental Measures. *Developmental Science*, 29(1), e70088. <https://doi.org/10.1111/desc.70088>
- Weaver, H., Saffran, J., & Arunachalam, S. (2026). Looking forward: Eye-gaze methods in vocabulary development research. *Trends in Cognitive Sciences*.
- Weaver, H., Zettersten, M., & Saffran, J. R. (2024). Becoming word meaning experts: Infants' processing of familiar words in the context of typical and atypical exemplars. *Child Development*, 95(5), e352–e372. <https://doi.org/10.1111/cdev.14120>
- Weisleder, A., & Fernald, A. (2013). Talking to children matters: Early language experience strengthens processing and builds vocabulary. *Psychological Science*, 24(11), 2143–2152. <https://doi.org/10.1177/0956797613488145>
- Yoon, E. J., Wu, Y. C., & Frank, M. C. (2015). Children's online processing of ad-hoc implicatures. *Proceedings of the Annual Meeting of the Cognitive Science Society*, 37.
- Yurovsky, D., & Frank, M. C. (2017). Beyond naïve cue combination: Salience and social cues in early word learning. *Developmental Science*, 20(2), e12349.

- Yurovsky, Daniel, Wade, A., & Frank, M. C. (2013). Online processing of speech and social information in early word learning. *Proceedings of the 35th Annual Meeting of the Cognitive Science Society*.
- Yurovsky, Daniel, Wade, A., Kraus, A., Gengoux, G., Hardan, A. Y., & Frank, M. C. (2024). *Developmental changes in the speed of social attention in early word learning*. PsyArXiv preprint. Retrieved from https://osf.io/preprints/psyarxiv/bcrj2_v1
- Zettersten, M., Yurovsky, D., Xu, T. L., Uner, S., Tsui, A. S. M., Schneider, R. M., . . . Frank, M. C. (2023). Peekbank: An open, large-scale repository for developmental eye-tracking data of children's word recognition. *Behavior Research Methods*, 55(5), 2485–2500. <https://doi.org/10.3758/s13428-022-01906-4>

Supplemental Information for Data-driven recommendations
for increasing reliability and validity in measures
of infant word recognition

Contents

1	Data description	1
2	Methods details	2
2.1	Criteria for inclusion	2
2.2	Details about ICC	3
3	Accuracy	4
4	Baseline correction	5
5	Reaction time	7
6	Exclusions	9
6.1	Accuracy trial exclusions	9
6.2	Accuracy exclusions by administration	10
6.3	Reaction time exclusions	12
7	Order effects	13
	References	15

1 Data description

Figure S1 shows the distribution of ages in the Peekbank data used in this paper.




Figure S1: Age distribution for trials and administrations, with the age bins used for analysis.

Figure S2 shows the distribution of AOIs for different datasets and different timepoints. Some datasets have less data at the start or end due to varying length recordings, or being a mix of multiple experiments with different recording windows. Most (but not all) datasets had data before the point of disambiguation and up to 3000ms. Datasets varied dramatically in overall size.




Figure S2: Distribution of AOIs across timepoints and datasets. Solid vertical line at point of disambiguation, Dashed lines at 4000ms, and dotted line at -2000ms and 2000ms.

2 Methods details

2.1 Criteria for inclusion

The “vanilla” status of a trial is coded in the Peekbank dataset. As part of the 2026.1 release, we implemented an updated set of “vanilla” criteria. Multiple coders systematically inspected the coding of each dataset to identify vanilla trials and ensured that the application of the vanilla criteria was consistent across all datasets. Cases where there was potential ambiguity were discussed and documented in the README associated with each dataset, to ensure that the criteria used were consistent and explainable. The criteria used were as follows:

Auditory Stimulus:

- familiar words/ target word is a real noun (no part words)
- target word is the first point of disambiguation

- no duplicated target words (i.e., repeating the target word multiple times was deemed non-vanilla)
- well-formed and grammatical carrier phrase (there were no restrictions on carrier vocabulary, and the target NP could appear anywhere in the sentence.)
- no informative structure prenominal or postnominal (informative verb or adjective; an uninformative verb or adjective is OK)
- no nonsense words anywhere (including the carrier phrase)
- no language/ speaker/ accent switching within the trial
- no intentional background noise or audio filtering
- no additional restrictions on noun difficulty

Visual stimuli:

- both target and distractor images are familiar objects (i.e., each could be valid familiar targets; non-prototypical objects are OK, but unnatural/ artificially modified exemplars are not)
- no novel objects are visible
- The target referent is the focal object of its image (dog on grass for dog is OK; apple on plate is OK; apple and orange for orange is not, because both the apple and orange are competing for focal attention).

2.2 Details about ICC

A number of different correlation metrics are all abbreviated as ICC. In this study, we use the ICC for *inter*-rater reliability for a “random factorial model without interaction” measuring consistency and looking for the reliability of an average of raters. We use the implementation in the “agreement” package (Girard 2019), but we are not aware of this specific combination being fully documented anywhere.

This ICC looks at the ratio between by-subject variance (by-subject variance + by-rater variance + error variance)

Here’s what this ICC means in comparison to other possible ICCs:

- It is an inter-rater reliability measure, meaning that it measures how well raters (or in our case items) correlate with each other as opposed to with themselves.
- It uses the “random factorial model” (sometimes called Type 2). This is for when all raters rate all subjects, but some raters may rate some subjects more than once. Both raters and subjects are assumed to be sampled from a larger population of interest.
- It is “without interaction”. This variant can be used when there are single ratings for a subject-rater combination (preventing the accurate determination of subject-rater interaction effects), or when subject-rater interactions are not of interest (and can be lumped into the overall error). This version is documented in Gwet (2014) (p. 256). This measures the ratio of by-subject variance / (by-subject variance + by-rater variance + error variance).
- It is a consistency measure rather than an agreement measure. Rather than measure the exact agreement between raters, it instead measures whether the slopes across raters for different subjects are the same. It does not penalize situations where one rater is consistently lower (i.e., a later-learned word). This omits the by-rater variance from the denominator. We are unaware of the combination of random-factorial model and consistency being explicitly documented, but this formula aligns with the discussion of consistency in Gwet (2014) (p. 274).
- It measures consistency on average across a pool of raters by dividing the error term by the number of raters.

The overall formula for ICC is then by-subject variance / (by-subject variance + (error variance / number of raters))

We use the package “agreement” because it, unlike other ICC packages, handles cases where there is missing or uneven data (by using Gwet’s formulations). This is essential because peekbank has uneven data, especially for cases involving RT or post-exclusion data.

We only use the ICC measure to compare across the same datasets with different processing methods applied. We do not interpret the ICC values against any external cutoffs of reliability.

3 Accuracy

Figure S3 shows a version of Figure 1 split by age bin categories. Older children have overall higher ICC reliability, but different ages do not show different regions of high reliability. Figure S4 shows the model estimates and 95% CIs for the effect of different windowing options (comparing to 600–4000 ms as the index condition). Negative numbers indicate that the value is lower in the given window than for the index condition; higher numbers indicate that the value is higher in the given window than in the index condition. Figure S5 shows the model coefficients and 95% CIs for linear mixed effect models run separately for each age bin. While the degrees of difference vary across ages, the overall patterns are the same.




Figure S3: A heat map showing the average ICC across datasets for different start and end points for the accuracy window, split by age bins. A region with the cross-age highest ICC (best reliability) is outlined.




Figure S4: Model results comparing the effects different accuracy windows to the index length of 600–4000 ms for different dependent variables.




Figure S5: Model results split out by age comparing the effects different accuracy windows to the index length of 600–4000 ms for different dependent variables.

4 Baseline correction

For baseline correction, we only considered datasets where the earliest time point was at or before -2000 ms, to ensure there was adequate per-onset data.

Figure S6 shows how baseline corrected measures compare to normal accuracy measures for a 300–1800 ms window, as an example of generalization. Figure S7 shows model results comparing different baseline correction options to a plain accuracy measure. Baseline correction is substantially worse. Figure S8 shows model results comparing baseline corrected accuracies to the index condition of normal accuracy for each of

the age bins. Across age bins, baseline correction is often worse and in no case better than normal accuracy.




Figure S6: Comparison between a baseline-corrected accuracy in red and accuracy in blue across datasets and measures. Both use a post-stimulus window of 300 - 1800ms, the baseline window was -2000 - 0ms.




Figure S7: Model results comparing the effect of different baseline windows for baseline correction to uncorrected accuracy. All use a post stimulus window of 600-4000 ms.




Figure S8: Model results by age group comparing the effect of different baseline windows for baseline correction to uncorrected accuracy. All use a post stimulus window of 600–4000 ms.

5 Reaction time

Figure S9 shows a heat map of ICC reliability across different RT conditions, split by age. Younger children are less reliable overall, but there are not strong age-specific trends in what measurements are most reliable.




Figure S9: Heat map of different parameter options for measuring RT and how they on average affected ICC.

Figure S10 shows models estimates for different model parameters compared to the index condition of untrimmed log land with 400 minimum RT. Because faster RT tends to correlate with higher language skill and thus higher CDI measures, for the bottom two panels, positive numbers indicate worse correlations in the tested condition compared to index. Figure S11 shows the reliability and validity of different reaction time parameters, split by age bin.

For thoroughness, we also considered other dimensions of RT parameters focused on the 400–4000 ms RT range.




Figure S10: Model results for RT. The index condition is a log Land RT from 400–4000 ms. All RTs are shown log in this figure.




Figure S11: Model results for RT. The index condition is a log Land RT from 400–4000 ms. All RTs are shown log in this figure.

In Figure S12, we consider land versus launch RT, log versus raw, and whether to exclude long shifts. While not formally mentioned, launch based RT is often restricted to RTs where the child's eyes make it to the

target in a reasonable amount of time (ex. 600 ms). This restriction on shift-length can also be applied to land based RT. As shown in the figure, these differences in calculation do not make significant differences.




Figure S12: Model results for different ways of computing RT on the 400–4000 ms window, focused on max-shift length and transformation.

6 Exclusions

6.1 Accuracy trial exclusions

Figure S13 shows effects of different trial-level exclusions on outcomes. Figure S14 shows age splits of the same.

Excessive requirements lead to an increase in ICC, but a decrease in other metrics. A requirement to look at the time of auditory stimulus onset just leads to data loss without reliability or validity gains.




Figure S13: Model results for possible trial-level exclusions, compared to no exclusions. Colors are overall rate of trial-level data retention.




Figure S14: Model results for possible trial-level exclusions for each age group, compared to no exclusions.

6.2 Accuracy exclusions by administration

Target number of total trials is not a peekbank field, so we guessed it by the largest number of trials any child in that dataset had, and adjusted for datasets where it was clear from the distributions there were

different numbers of vanilla trials at different ages or administrations.

Overall model results are shown for trial count exclusions in Figure S15 and for trial fraction exclusions in Figure S16. Age splits are shown in Figure S17 and Figure S18.

More trials numerically improves test-retest reliability (at least when there are enough trials), but there are trade-offs against number of retained administrations.




Figure S15: Model results for possible trial count based exclusions, compared to no minimum trial count. Colors indicate fraction of administrations retained.




Figure S16: Model results for possible fraction of trial based exclusions, compared to no minimum. Colors indicate fraction of administrations retained.




Figure S17: Model results for different age groups for possible trial count based exclusions, compared to no minimum trial count.




Figure S18: Model results for different age groups for possible fraction of trial based exclusions, compared to no minimum.

6.3 Reaction time exclusions

Excluding children who have fewer valid RT trials may increase test-retest reliability, and has unclear relationships with validity. Overall model results are shown in Figure S19 and model results by age group are down in Figure S20. Minimum trial requirements do help for test-retest reliability, and CDI comprehension, at the cost of reducing the number of included administrations.




Figure S19: Effects of imposing minimum RT trial requirements, compared to no trial minimum, colored by what fraction of administrations overall are retained.




Figure S20: Effects of imposing minimum RT trial requirements, compared to no trial minimum, split by age bin.

7 Order effects

One potential concern, especially when children see a number of trials, is that the repetition of target and distractor images (or just the target nouns) may lead to order effects or preferences for familiarity, i.e., the image that was named last. It is unclear how these effects might manifest, so we looked for possible order effects in a combinatorial fashion.

We considered 5 dependent variables:

- “early looking” on the -2000 to 0 pre-window
- “late looking” on the 0 to 600 window (between onset of stimulus name and when we start an accuracy window)
- accuracy on 600–4000 ms
- RT on 400–4000 ms
- whether there is a valid RT (if children’s looks at 0 are disproportionate, this could reduce the trials where RT is possible)

We also considered a set of independent variables across different models:

- trial order: for stimuli seen for the first time, how does position in order change viewing behavior?
- seen previously (as distractor) – for stimuli, how does being seen for the first time as the target compare to being the target when seen once previously as the distractor? (We restrict to pairs of stimuli that are always each other’s distractors.)
- seen previously (as distractor) x trial order
- seen previously (as target) – for stimuli, how does being seen for the first time as the target compare to being the target when seen once previously as the target? (We restrict to pairs of stimuli that are always each other’s distractors)
- seen previously (as target) x trial order
- seen previously (as either) – for datasets where the same items could be seen previously as either the target or distractor, we do a 3-way comparison between first time viewing and one previous viewing as either target or distractor.
- seen previously (as either) x trial order

We fully crossed the 5 dependent variables with the 7 model predictor structures.

The coefficients from all these models are shown in Figure S21. Effects related to trial order were multiplied by 10 for the effect of a difference in 10 trials for viewability. The effects of having a target image that has been seen 1 time previously (as distractor or target), versus a target image that is being seen for the first time are small, around 1% differences in looks. These effects may well interact with trial order, although again, effects and interactions, if they exist, are small. Log RT has a larger scale than the other measures that are all on 0–1. Switches to target may be a bit faster if the target was previously seen, but only when interaction with trial order is considered in the model. If anything, there is more looking to previous distractors on the accuracy window. At least across datasets, we cannot detect reliable effects of previously seen items or trial order on any of the dependent variables.




Figure S21: Fitted coefficient estimates from a series of models investigating potential order and repetition effects. For scaling, trial order effects are shown for every 10 trials.

References

Girard, Jeffrey. 2019. *Agreement: R Package for the Tidy Calculation of Inter-Rater Reliability*. <https://github.com/jmgirard/agreement>.

Gwet, Kilem Li. 2014. *Handbook of Inter-Rater Reliability: The Definitive Guide to Measuring the Extent of Agreement Among Raters*. 4th ed. Advanced Analytics, LLC.